

# Is human oversight of this system doing real work?

Work through this against one real system you are responsible for — not a class of systems, and not the policy that covers them. The answers are only worth the evidence written beside them.

Mubien Ahsan · [mubienahsan.com/research/autonomy-governance](https://mubienahsan.com/research/autonomy-governance) · Sources last checked 20 September 2026

## How to use this

- 1. Name the system and the person who is actually accountable for its decisions, not the team that runs it.
- 2. Place it on the autonomy scale in part two. Place it by what it can do, not by what the documentation says it is for.
- 3. Answer the five oversight questions in part three, writing the evidence next to each. “We would notice” is not evidence; a case where someone did notice is.
- 4. Total the score and read the band in part four.
- 5. Check the governance dimensions in part five against the intensity the level calls for, and record the gaps.

This is a structured way to have an argument, not a control test or an assurance procedure. It produces a view and a date to revisit it. It does not produce a rating, a certification or a compliance conclusion.

## Part 1 · The system

|                                                                |                 |
|----------------------------------------------------------------|-----------------|
| SYSTEM OR AGENT NAME                                           | ASSESSMENT DATE |
| ACCOUNTABLE PERSON (NAMED, NOT A TEAM)                         | ASSESSED BY     |
| WHAT IT IS FOR, IN ONE LINE                                    |                 |
| WHAT IT IS EXPLICITLY NOT FOR                                  |                 |
| WHAT IT CAN REACH: DATA, TOOLS, SYSTEMS, COUNTERPARTIES        |                 |
| WHAT IT CAN DO WITHOUT A PERSON APPROVING THAT SPECIFIC ACTION |                 |

## Part 2 · Where it sits on the autonomy scale

---

Tick the highest level the system can reach in practice. If people disagree, that disagreement is the finding — record it rather than resolving it by picking the lower one.

☐ **Level 1 — Assistive**

Produces information, analysis, a recommendation or a draft. It does not act on the world, and nothing it writes takes effect until a person does something with it.

---

☐ **Level 2 — Delegated task execution**

Carries out a clearly bounded task — a defined input, a defined output, narrow limits. It can touch systems, but the consequential step waits for approval.

---

☐ **Level 3 — Conditional autonomy**

Plans an approach, chooses tools, and takes a sequence of actions without approval for each one — as long as it stays inside defined limits.

---

☐ **Level 4 — High autonomy**

Works towards a longer-running objective, adapts its strategy as it goes, and may coordinate several systems or other agents with limited intervention.

---

☐ **Level 5 — Transformative autonomy (scenario, not a claim that this exists)**

In this scenario the system outperforms people across many relevant tasks, can accelerate parts of AI research, or can improve important aspects of its own operation.

---

WHAT WOULD MOVE IT UP A LEVEL, AND IS ANYTHING ALREADY MOVING IT?

---

## Part 3 · The Meaningful Human Oversight Test

---

Five separable conditions. A system can satisfy four and still fail, because the one it misses is the one that mattered. Score yes, demonstrably = 2, partly, or untested = 1, no = 0.

### 1. Understand — Can the responsible person understand the objective, the boundaries, the relevant inputs and the proposed action?

Everything else depends on this one. You cannot challenge, verify or sensibly intervene in something you do not follow.

Failure looks like: The reviewer can describe what the system produced but not why, and fills the gap with an assumption about how it probably works.

☐ Yes, demonstrably (2)    ☐ Partly, or untested (1)    ☐ No (0)

EVIDENCE FOR THAT ANSWER — WHAT SPECIFICALLY, AND WHEN WAS IT LAST TRUE?

---

IF PARTLY OR NO: WHAT WOULD HAVE TO CHANGE

---

---

### 2. Challenge — Can that person question the recommendation and reach a different conclusion, without simply deferring to it?

Deference is the quiet failure. A reviewer who always agrees is indistinguishable, in the record, from no reviewer at all.

Failure looks like: Overriding is possible in principle and never happens in practice, and nobody has asked why.

☐ Yes, demonstrably (2)    ☐ Partly, or untested (1)    ☐ No (0)

EVIDENCE FOR THAT ANSWER — WHAT SPECIFICALLY, AND WHEN WAS IT LAST TRUE?

---

IF PARTLY OR NO: WHAT WOULD HAVE TO CHANGE

---

---

**3. Verify — Can they independently test the important facts, assumptions or outputs before harm occurs?**

Verification is not review. Review asks whether the output looks right; verification checks it against something the system did not supply.

Failure looks like: The only evidence available to the reviewer is the evidence the system produced, so agreement is the only reachable outcome.

☐ Yes, demonstrably (2)    ☐ Partly, or untested (1)    ☐ No (0)

EVIDENCE FOR THAT ANSWER — WHAT SPECIFICALLY, AND WHEN WAS IT LAST TRUE?

IF PARTLY OR NO: WHAT WOULD HAVE TO CHANGE

**4. Intervene — Can they interrupt or change the behaviour while the decision is still reversible?**

Timing is the whole control. The ability to intervene after the money moves is not intervention, it is reporting.

Failure looks like: Intervention exists and is slower than the process it is meant to interrupt, which nobody has measured.

☐ Yes, demonstrably (2)    ☐ Partly, or untested (1)    ☐ No (0)

EVIDENCE FOR THAT ANSWER — WHAT SPECIFICALLY, AND WHEN WAS IT LAST TRUE?

IF PARTLY OR NO: WHAT WOULD HAVE TO CHANGE

**5. Stop — Can they reliably end the system's authority, access and ability to act across every connected tool?**

Stopping the interface is not stopping the system. Credentials, tools, queued work and delegated agents each need to be reached.

Failure looks like: A documented shutdown procedure that has never been exercised, so its real duration and its gaps are both unknown.

☐ Yes, demonstrably (2)    ☐ Partly, or untested (1)    ☐ No (0)

EVIDENCE FOR THAT ANSWER — WHAT SPECIFICALLY, AND WHEN WAS IT LAST TRUE?

IF PARTLY OR NO: WHAT WOULD HAVE TO CHANGE

Total score, out of 10 \_\_\_\_\_

## Part 4 · Reading the result

---

### 9–10 · Meaningful oversight appears possible

On your own answers, the person responsible could understand, challenge, verify, intervene and stop. That is the condition under which human review is doing real work.

Next: Keep it true. Re-run this when the model, the tools, the permissions or the volume of decisions change — any of those can move the answer without anyone deciding to.

---

### 7–8 · Oversight is fragile

Most of the conditions hold, but at least one is weak or untested. Fragile oversight tends to look fine until the day it is needed, because the gap only shows under pressure.

Next: Find the weakest of the five and make it real rather than nominal. An untested stop and an unmeasured intervention window are the two most common.

---

### 4–6 · Oversight risks becoming ceremonial

There is a person in the workflow, and on these answers they may not be able to reach a different conclusion from the system's. The approval records a decision without evidencing a check.

Next: Either strengthen the reviewer's position — information, time, independence, authority — or stop relying on human review as the primary control and put the weight on constraints that do not depend on it.

---

### 0–3 · Human review alone is not an adequate control

On these answers, the person cannot verify the system in the time available. Keeping an approval step here does not reduce the risk; it mainly documents who was present.

Next: Reduce what the system may do without a constraint that holds independently of a person: narrower permissions, enforced limits, reversible actions, smaller blast radius. Then re-test.

---

BAND REACHED

DATE TO REASSESS

THE WEAKEST OF THE FIVE, AND WHO OWNS CLOSING IT

---

Re-run this whenever the model, the tools, the permissions, the operating context or the volume of decisions changes. Any of those can move the answer without anyone deciding that it should.

## Part 5 · Governance dimensions against the level

For the level you ticked in part two, the framework calls for the intensity shown below. Record what is actually in place, the evidence it operated, and the gap. A control that exists and has never been exercised belongs in the gap column.

| Dimension                                                                                                | L1 | L2 | L3 | L4 | L5 | In place / evidence | Gap and owner |
|----------------------------------------------------------------------------------------------------------|----|----|----|----|----|---------------------|---------------|
| <b>Purpose and permitted use</b><br>What is this system for, and what is it explicitly not for?          | 1  | 2  | 4  | 5  | 5  |                     |               |
| <b>Authority and permissions</b><br>What may it actually reach and do, enforced outside the model?       | 1  | 3  | 4  | 5  | 5  |                     |               |
| <b>Human oversight</b><br>Who reviews what, with what information, and can they say no?                  | 2  | 3  | 3  | 4  | 5  |                     |               |
| <b>Testing and evaluation</b><br>What was tested, against what expected behaviour, and by whom?          | 1  | 3  | 4  | 5  | 5  |                     |               |
| <b>Monitoring and evidence</b><br>What is watched in production, and what would the record show?         | 1  | 3  | 4  | 5  | 5  |                     |               |
| <b>Intervention and containment</b><br>How is it stopped, how fast, and has that been exercised?         | 1  | 2  | 4  | 5  | 5  |                     |               |
| <b>Accountability and escalation</b><br>Who owns the outcome, and where does a serious problem go?       | 2  | 3  | 4  | 5  | 5  |                     |               |
| <b>Independent assurance</b><br>Who checks this who did not build it and does not report to whoever did? | 1  | 2  | 3  | 5  | 5  |                     |               |

1 = light touch, 5 = the heaviest arrangement this framework describes. Light · Moderate · Substantial · Heavy · Maximal. These are relative weights for comparing levels, not measurements and not a maturity score.

## Part 6 · What this did not test

---

Recording the limits keeps the document honest when it is read later by someone who was not in the room. This template does not test whether the model is accurate, whether it is robust to prompt injection, whether its reasoning can be interpreted, or whether a shutdown works — those need their own work, and none of them is a solved problem.

### KNOWN LIMITATIONS OF THIS ASSESSMENT

---

### WHAT WAS ASSUMED RATHER THAN VERIFIED

---

### WHO DISAGREED, AND WITH WHAT

---

---

Written in a personal capacity. It is not the position of any employer, is not derived from any employer’s methodology or client work, and is not legal, regulatory or professional advice. Nothing here is a certification, an audit opinion or a compliance conclusion.

Voluntary frameworks and binding law are distinguished throughout and should not be read as interchangeable. This area moves quickly; check the sources against their current versions before relying on them.

The interactive version of this test, the full framework and the sources behind it are at [mubienahsan.com/research/autonomy-governance](https://mubienahsan.com/research/autonomy-governance).

© 2026 Mubien Ahsan. Free to use and share within your organisation with attribution.